
Beaker Documentation

Release 1.5.2

Ben Bangert, Mike Bayer

March 01, 2010

CONTENTS

1	Configuration	1
1.1	Options For Sessions and Caching	2
1.2	Session Options	2
1.3	Cache Options	3
2	Sessions	5
2.1	About	5
2.2	Using	5
2.3	Saving	6
2.4	Deleting	7
2.5	Cookie Domain and Path	7
2.6	Cookie-Based	8
3	Caching	9
3.1	About	9
3.2	Creating the CacheManager Instance	9
3.3	Programmatic API	10
3.4	Decorator API	10
3.5	Cache Regions	12
4	Changes in Beaker	13
4.1	Release 1.5.2 (3/1/2010)	13
4.2	Release 1.5.1 (12/17/2009)	13
4.3	Release 1.5 (11/23/2009)	13
4.4	Release 1.4.2 (9/25/2009)	14
4.5	Release 1.4.1 (9/10/2009)	14
4.6	Release 1.4 (7/24/2009)	14
4.7	Release 1.3.1 (5/5/2009)	14
4.8	Release 1.3 (4/6/2009)	15
4.9	Release 1.2.3 (3/2/2009)	15
4.10	Release 1.2.2 (2/14/2009)	15
4.11	Release 1.2.1 (2/09/2009)	15
4.12	Release 1.2 (1/22/2009)	16
4.13	Release 1.1.3 (12/29/2008)	16
4.14	Release 1.1.2 (11/24/2008)	16
4.15	Release 1.1.1 (11/24/2008)	16
4.16	Release 1.1 (11/16/2008)	16
4.17	Release 1.0.3 (10/14/2008)	17
4.18	Release 1.0.2 (9/22/2008)	17

4.19	Release 1.0.1 (8/19/2008)	17
4.20	Release 1.0 (8/13/2008)	17
4.21	Release 0.9.5 (6/19/2008)	17
4.22	Release 0.9.4 (4/13/2008)	17
4.23	Release 0.9.3 (2/28/2008)	18
4.24	Release 0.9.2 (2/13/2008)	18
4.25	Release 0.9.1 (2/4/2008)	18
4.26	Release 0.9 (12/17/07)	18
4.27	Release 0.8.1 (11/15/07)	19
4.28	Release 0.8 (10/17/07)	19
4.29	Release 0.7.5 (08/18/07)	19
4.30	Release 0.7.4 (07/09/07)	19
4.31	Release 0.7.3 (06/08/07)	19
4.32	Release 0.7.2 (05/19/07)	20
4.33	Release 0.7.1 05/18/07)	20
4.34	Release 0.7 (05/18/07)	20
4.35	Release 0.6.3 (03/18/2007)	20
5	Indices and tables	21
5.1	Module Listing	21

CONFIGURATION

Beaker can be configured several different ways, depending on how it's used. The most recommended style is to use a dictionary of preferences that are to be passed to either the `SessionMiddleware` or the `CacheManager`.

Since both Beaker's sessions and caching use the same back-end container storage system, there's some options that are applicable to both of them in addition to session and cache specific configuration.

Most options can be specified as a string (necessary to config options that are setup in INI files), and will be coerced to the appropriate value. Only datetime's and timedelta's cannot be coerced and must be the actual objects.

Frameworks using Beaker usually allow both caching and sessions to be configured in the same spot, Beaker assumes this condition as well and requires options for caching and sessions to be prefixed appropriately.

For example, to configure the `cookie_expires` option for Beaker sessions below, an appropriate entry in a `Pylons` INI file would be:

```
# Ironically, setting cookie_expires = true causes Beaker to omit the
# expires= field from the Set-Cookie: header, signaling the cookie
# should be discarded when the browser closes.
beaker.session.cookie_expires = true
```

Note: When using the options in a framework like `Pylons` or `TurboGears2`, these options must be prefixed by `beaker.`, for example in a `Pylons` INI file:

```
beaker.session.data_dir = %(here)s/data/sessions/data
beaker.session.lock_dir = %(here)s/data/sessions/lock
```

Or when using stand-alone with the `SessionMiddleware`:

```
from beaker.middleware import SessionMiddleware

session_opts = {
    'session.cookie_expires' : True
}

app = SomeWSGIAPP()
app = SessionMiddleware(app, session_opts)
```

Or when using the `CacheManager`:

```
from beaker.cache import CacheManager
from beaker.util import parse_cache_config_options

cache_opts = {
    'cache.type' : 'file',
    'cache.data_dir' : '/tmp/cache/data',
    'cache.lock_dir' : '/tmp/cache/lock'
}

cache = CacheManager(**parse_cache_config_options(cache_opts))
```

Note: When using the `CacheManager` directly, all dict options must be run through the `beaker.util.parse_cache_config_options()` function to ensure they're valid and of the appropriate type.

1.1 Options For Sessions and Caching

data_dir (optional, string) Used with any back-end that stores its data in physical files, such as the dbm or file-based back-ends. This path should be an absolute path to the directory that stores the files.

lock_dir (required, string) Used with every back-end, to coordinate locking. With caching, this lock file is used to ensure that multiple processes/threads aren't attempting to re-create the same value at the same time (*The Dog-Pile Effect*)

type (required, string) The name of the back-end to use for storing the sessions or cache objects.

Available back-ends supplied with Beaker: `file`, `dbm`, `memory`, `ext:memcached`, `ext:database`, `ext:google`

For sessions, the additional type of `cookie` is available which will store all the session data in the cookie itself. As such, size limitations apply (4096 bytes).

Some of these back-ends require the `url` option as listed below.

url (optional, string) URL is specific to use of either `ext:memcached` or `ext:database`. When using one of those types, this option is **required**.

When used with `ext:memcached`, this should be either a single, or semi-colon separated list of memcached servers:

```
session_opts = {
    'session.type' : 'ext:memcached',
    'session.url' : '127.0.0.1:11211',
}
```

When used with `ext:database`, this should be a valid [SQLAlchemy](#) database string.

1.2 Session Options

The Session handling takes a variety of additional options relevant to how it stores session id's in cookies, and when using the optional encryption.

auto (optional, bool) When set to `True`, the session will save itself anytime it is accessed during a request, negating the need to issue the `save()` method.

Defaults to False.

cookie_expires (optional, bool, datetime, timedelta) Determines when the cookie used to track the client-side of the session will expire. When set to a boolean value, it will either expire at the end of the browsers session, or never expire.

Setting to a datetime forces a hard ending time for the session (generally used for setting a session to a far off date).

Defaults to never expiring.

cookie_domain (optional, string) What domain the cookie should be set to. When using sub-domains, this should be set to the main domain the cookie should be valid for. For example, if a cookie should be valid under `www.nowhere.com` **and** `files.nowhere.com` then it should be set to `.nowhere.com`.

Defaults to the current domain in its entirety.

Alternatively, the domain can be set dynamically on the session by calling, see [Session Attributes / Keys](#).

key (required, string) Name of the cookie key used to save the session under.

secret (required, string) Used with the HMAC to ensure session integrity. This value should ideally be a randomly generated string.

When using in a cluster environment, the secret must be the same on every machine.

secure (optional, bool) Whether or not the session cookie should be marked as secure. When marked as secure, browsers are instructed to not send the cookie over anything other than an SSL connection.

timeout (optional, integer) Seconds until the session is considered invalid, after which it will be ignored and invalidated. This number is based on the time since the session was last accessed, not from when the session was created.

Defaults to never expiring.

1.2.1 Encryption Options

These options should then be used *instead* of the `secret` option listed above.

encrypt_key (required, string) Encryption key to use for the AES cipher. This should be a fairly long randomly generated string.

validate_key (required, string) Validation key used to sign the AES encrypted data.

Note: You may need to install additional libraries to use Beaker's cookie-based session encryption. See the [Encryption](#) section for more information.

1.3 Cache Options

For caching, options may be directly specified on a per-use basis with the `cache()` decorator, with the rest of these options used as fallback should one of them not be specified in the call.

Only the `lock_dir` option is strictly required, unless using the file-based back-ends as noted with the sessions.

enabled (optional, bool) Quick toggle to disable or enable caching across an entire application.

This should generally be used when testing an application or in development when caching should be ignored.

Defaults to True.

expire (optional, integer) Seconds until the cache is considered old and a new value is created.

1.3.1 Cache Region Options

Starting in Beaker 1.3, cache regions are now supported. These can be thought of as bundles of configuration options to apply, rather than specifying the type and expiration on a per-usage basis.

regions (optional, list, tuple) Names of the regions that are to be configured.

For each region, all of the other cache options are valid and will be read out of the cache options for that key. Options that are not listed under a region will be used globally in the cache unless a region specifies a different value.

For example, to specify two batches of options, one called `long-term`, and one called `short-term`:

```
cache_opts = {
    'cache.data_dir': '/tmp/cache/data',
    'cache.lock_dir': '/tmp/cache/lock',
    'cache.regions': 'short_term, long_term',
    'cache.short_term.type': 'ext:memcached',
    'cache.short_term.url': '127.0.0.1:11211',
    'cache.short_term.expire': '3600',
    'cache.long_term.type': 'file',
    'cache.long_term.expire': '86400',
```


SESSIONS

2.1 About

Sessions provide a place to persist data in web applications, Beaker's session system simplifies session implementation details by providing WSGI middleware that handles them.

All cookies are signed with an HMAC signature to prevent tampering by the client.

2.1.1 Lazy-Loading

Only when a session object is actually accessed will the session be loaded from the file-system, preventing performance hits on pages that don't use the session.

2.2 Using

The session object provided by Beaker's `SessionMiddleware` implements a dict-style interface with a few additional object methods. Once the `SessionMiddleware` is in place, a session object will be made available as `beaker.session` in the WSGI environ.

Getting data out of the session:

```
myvar = session[ ' somekey ' ]
```

Testing for a value:

```
logged_in = ' user_id ' in session
```

Adding data to the session:

```
session[ ' name ' ] = ' Fred Smith '
```

Complete example using a basic WSGI app with sessions:

```
from beaker.middleware import SessionMiddleware

def simple_app(environ, start_response):
    # Get the session object from the environ
```

```
session = environ[ ' beaker.session ' ]

# Check to see if a value is in the session
if ' logged_in ' in session:
    user = True
else:
    user = False

# Set some other session variable
session[ ' user_id ' ] = 10

start_response( ' 200 OK ' , [ ( ' Content-type ' , ' text/plain ' ) ])
return [ ' User is logged in: %s ' % user]

# Configure the SessionMiddleware
session_opts = {
    ' session.type ' : ' file ' ,
    ' session.cookie_expires ' : True,
}
wsgi_app = SessionMiddleware(simple_app, session_opts)
```

Note: This example does **not** actually save the session for the next request. Adding the `save()` call explained below is required, or having the session set to auto-save.

2.2.1 Session Attributes / Keys

Sessions have several special attributes that can be used as needed by an application.

- `id` - Unique 40 char SHA-generated session ID
- `last_accessed` - The last time the session was accessed before the current access, will be `None` if the session was just made

There's several special session keys populated as well:

- `_accessed_time` - Current accessed time of the session, when it was loaded
- `_creation_time` - When the session was created

2.3 Saving

Sessions can be saved using the `save()` method on the session object:

```
session.save()
```

Warning: Beaker relies on Python's pickle module to pickle data objects for storage in the session. Objects that cannot be pickled should **not** be stored in the session.

This flags a session to be saved, and it will be stored on the chosen back-end at the end of the request.

If it's necessary to immediately save the session to the back-end, the `persist()` method should be used:

```
session.persist()
```

This is not usually the case however, as a session generally should not be saved should something catastrophic happen during a request.

2.3.1 Auto-save

Saves can be done automatically by setting the `auto` configuration option for sessions. When set, calling the `save()` method is no longer required, and the session will be saved automatically anytime its accessed during a request.

2.4 Deleting

Calling the `delete()` method deletes the session from the back-end storage and sends an expiration on the cookie requesting the browser to clear it:

```
session.delete()
```

This should be used at the end of a request when the session should be deleted and will not be used further in the request.

If a session should be invalidated, and a new session created and used during the request, the `invalidate()` method should be used:

```
session.invalidate()
```

2.4.1 Removing Expired/Old Sessions

Beaker does **not** automatically delete expired or old cookies on any of its back-ends. This task is left up to the developer based on how sessions are being used, and on what back-end.

The database backend records the last accessed time as a column in the database so a script could be run to delete session rows in the database that haven't been used in a long time.

When using the file-based sessions, a script could run to remove files that haven't been touched in a long time, for example (in the session's data dir):

```
find . -mtime +3 -exec rm {} \;
```

2.5 Cookie Domain and Path

In addition to setting a default cookie domain with the *cookie domain setting*, the cookie's domain and path can be set dynamically for a session with the `domain` and `path` properties.

These settings will persist as long as the cookie exists, or until changed.

Example:

```
# Setting the session's cookie domain and path
session.domain = ' .domain.com '
session.path = ' /admin '
```

2.6 Cookie-Based

Session can be stored purely on the client-side using cookie-based sessions. This option can be turned on by setting the session type to `cookie`.

Using cookie-based session carries the limitation of how large a cookie can be (generally 4096 bytes). An exception will be thrown should a session get too large to fit in a cookie, so using cookie-based session should be done carefully and only small bits of data should be stored in them (the users login name, admin status, etc.).

Large cookies can slow down page-loads as they increase latency to every page request since the cookie is sent for every request under that domain. Static content such as images and Javascript should be served off a domain that the cookie is not valid for to prevent this.

Cookie-based sessions scale easily in a clustered environment as there's no need for a shared storage system when different servers handle the same session.

2.6.1 Encryption

In the event that the cookie-based sessions should also be encrypted to prevent the user from being able to decode the data (in addition to not being able to tamper with it), Beaker can use 256-bit AES encryption to secure the contents of the cookie.

Depending on the Python implementation used, Beaker may require an additional library to provide AES encryption.

On CPython (the regular Python), the `pycryptopp` library or `PyCrypto` library is required.

On Jython, no additional packages are required, but at least on the Sun JRE, the size of the encryption key is by default limited to 128 bits, which causes generated sessions to be incompatible with those generated in CPython, and vice versa. To overcome this limitation, you need to install the unlimited strength jurisdiction policy files from Sun:

- [Policy files for Java 5](#)
- [Policy files for Java 6](#)

CACHING

3.1 About

Beaker's caching system was originally based off the Perl `Cache::Cache` module, which was ported for use in **Myghty**. Beaker was then extracted from this code, and has been substantially rewritten and modernized.

Several concepts still exist from this origin though. Beaker's caching (and its sessions, though its behind the scenes) utilize the concept of *NamespaceManager*, and *Container* objects to handle storing cached data.

Each back-end utilizes a customized version of each of these objects to handle storing data appropriately depending on the type of the back-end.

The `CacheManager` is responsible for getting the appropriate *NamespaceManager*, which then stores the cached values. Each namespace corresponds to a single `thing` that should be cached. Usually a single `thing` to be cached might vary slightly depending on parameters, for example a template might need several different copies of itself stored depending on whether a user is logged in or not. Each one of these copies is then `keyed` under the *NamespaceManager* and stored in a *Container*.

There are three schemes for using Beaker's caching, the first and more traditional style is the programmatic API. This exposes the namespace's and retrieves a `Cache` object that handles storing keyed values in a *NamespaceManager* with *Container* objects.

The more elegant system, introduced in Beaker 1.3, is to use the *cache decorators*, these also support the use of *Cache Regions*.

Introduced in Beaker 1.5 is a more flexible `cache_region()` decorator capable of decorating functions for use with Beaker's *Cache Regions* **before** Beaker has been configured. This makes it possible to easily use Beaker's region caching decorator on functions in the module level.

3.2 Creating the CacheManager Instance

Before using Beaker's caching, an instance of the `CacheManager` class should be created. All of the examples below assume that it has already been created.

Creating the cache instance:

```
from beaker.cache import CacheManager
from beaker.util import parse_cache_config_options

cache_opts = {
    'cache.type' : 'file',
    'cache.data_dir' : '/tmp/cache/data',
```

```
' cache.lock_dir ' : ' /tmp/cache/lock '
}

cache = CacheManager(**parse_cache_config_options(cache_opts))
```

Additional configuration options are documented in the *Configuration* section of the Beaker docs.

3.3 Programmatic API

To store data for a cache value, first, a `NamespaceManager` has to be retrieved to manage the keys for a thing to be cached:

```
# Assuming that cache is an already created CacheManager instance
tmpl_cache = cache.get_cache(' mytemplate.html ', type=' dbm ', expire=3600)
```

Note: In addition to the defaults supplied to the `CacheManager` instance, any of the Cache options can be changed on a per-namespace basis, as this example demonstrates by setting a `type`, and `expire` option.

Individual values should be stored using a creation function, which will be called anytime the cache has expired or a new copy needs to be made. The creation function must not accept any arguments as it won't be called with any. Options affecting the created value can be passed in by using closure scope on the creation function:

```
search_param = ' gophers '

def get_results():
    # do something to retrieve data
    data = get_data(search_param)
    return data

# Cache this function, based on the search_param, using the tmpl_cache
# instance from the prior example
results = tmpl_cache.get(key=search_param, createfunc=get_results)
```

3.3.1 Invalidating

All of the values for a particular namespace can be removed by calling the `clear()` method:

```
tmpl_cache.clear()
```

Note that this only clears the key's in the namespace that this particular Cache instance is aware of. Therefore it's recommended to manually clear out specific keys in a cache namespace that should be removed:

```
tmpl_cache.remove_value(key=search_param)
```

3.4 Decorator API

When using the decorator API, a namespace does not need to be specified and will instead be created for you with the name of the module + the name of the function that will have its output cached.

Since its possible that multiple functions in the same module might have the same name, additional arguments can be provided to the decorators that will be used in the namespace to prevent multiple functions from caching their values in the same location.

For example:

```
# Assuming that cache is an already created CacheManager instance
@cache.cache(' my_search_func ', expire=3600)
def get_results(search_param):
    # do something to retrieve data
    data = get_data(search_param)
    return data

results = get_results(' gophers ')
```

The non-keyword arguments to the `cache()` method are the additional ones used to ensure this function's cache results won't clash with another function in this module called `get_results`.

The cache expire argument is specified as a keyword argument. Other valid arguments to the `get_cache()` method such as `type` can also be passed in.

When using the decorator, the function to cache can have arguments, which will be used as the key was in the *Programmatic API* for the data generated.

Warning: These arguments can **not** be keyword arguments.

3.4.1 Invalidating

Since the `cache()` decorator hides the namespace used, manually removing the key requires the use of the `invalidate()` function. To invalidate the 'gophers' result that the prior example referred to:

```
cache.invalidate(get_results, ' my_search_func ', ' gophers ')
```

If however, a type was specified for the cached function, the type must also be given to the `invalidate()` function so that it can remove the value from the appropriate back-end.

Example:

```
# Assuming that cache is an already created CacheManager instance
@cache.cache(' my_search_func ', type=" file ", expire=3600)
def get_results(search_param):
    # do something to retrieve data
    data = get_data(search_param)
    return data

cache.invalidate(get_results, ' my_search_func ', ' gophers ', type=" file ")
```

Note: Both the arguments used to specify the additional namespace info to the cache decorator **and** the arguments sent to the function need to be given to the `region_invalidate()` function so that it can properly locate the namespace and cache key to remove.

3.5 Cache Regions

Rather than having to specify the expiration, or toggle the type used for caching different functions, commonly used cache parameters can be defined as *Cache Regions*. These user-defined regions than may be used with the `region()` decorator rather than passing the configuration.

This can be useful if there are a few common cache schemes used by an application that should be setup in a single place then used as appropriate throughout the application.

Setting up cache region's is documented in the *cache region options* section in *Configuration*.

Assuming a `long_term` and `short_term` region were setup, the `region()` decorator can be used:

```
@cache.region(' short_term ', ' my_search_func ' )
def get_results(search_param):
    # do something to retrieve data
    data = get_data(search_param)
    return data

results = get_results(' gophers ' )
```

Or using the `cache_region()` decorator:

```
@cache_region(' short_term ', ' my_search_func ' )
def get_results(search_param):
    # do something to retrieve data
    data = get_data(search_param)
    return data

results = get_results(' gophers ' )
```

The only difference with the `cache_region()` decorator is that the cache does not need to be configured when its used. This allows one to decorate functions in a module before the Beaker cache is configured.

3.5.1 Invalidating

Since the `region()` decorator hides the namespace used, manually removing the key requires the use of the `region_invalidate()` function. To invalidate the 'gophers' result that the prior example referred to:

```
cache.region_invalidate(get_results, None, ' my_search_func ', ' gophers ' )
```

Or when using the `cache_region()` decorator, the `beaker.cache.region_invalidate()` function should be used:

```
region_invalidate(get_results, None, ' my_search_func ', ' gophers ' )
```

Note: Both the arguments used to specify the additional namespace info to the cache decorator **and** the arguments sent to the function need to be given to the `region_invalidate()` function so that it can properly locate the namespace and cache key to remove.

CHANGES IN BEAKER

4.1 Release 1.5.2 (3/1/2010)

- pkg_resources scanning for additional Beaker back-ends gracefully handles situations where its not present (GAE). Fixes #36.
- Avoid timing attacks on hash comparison.
- Provide abstract base for MemoryNamespaceManager that deals with “dictionaries”.
- Added tests for invalidating cache, and fixed bug with function cache when no args are present.
- The SQLAlchemy backends require SQLAlchemy 0.4 or greater (0.6 recommended).
- Rudimental Python 3 support is now available. Simply use Python 3 with Distribute and “python setup.py install” to run 2to3 automatically, or manually run 2to3 on “beaker” and “tests” to convert to a Python 3 version.
- Added support for PyCrypto module to encrypted session, etc. in addition to the existing pycryptopp support.

4.2 Release 1.5.1 (12/17/2009)

- Fix cache namespacing.

4.3 Release 1.5 (11/23/2009)

- Update memcached to default to using pylibmc when available.
- Fix bug when cache value doesn't exist causing has_key to throw an exception rather than return False. Fixes #24.
- Fix bug where getpid under GAE is used improperly to assume it should be a non-string. Fixes #22.
- Add cache_region decorator that works *before* configuration of the cache regions have been completed for use in module-level decorations.
- Fix bug where has_value sees the value before its removed.
- Improved accuracy of “dogpile” checker by removing dependency on “self” attributes, which seem to be slightly unreliable in highly concurrent scenarios.

4.4 Release 1.4.2 (9/25/2009)

- Fix bug where memcached may yank a value after the `has_value` but before the value can be fetched.
- Fix properties for setting the path. Fixes #15.
- Fix the 'TypeError: argument must be an int, or have a `fileno()` method' error sporadically emitted by `FileSynchronizer` under moderate load.

4.5 Release 1.4.1 (9/10/2009)

- Fix verification of options to throw an error if a beaker param is an empty string.
- Add `CacheManager.invalidate` function to easily invalidate cache spaces created by the use of the cache decorator.
- Add `CacheManager.region_invalidate` function to easily invalidate cache spaces created by the use of the `cache_region` decorator.
- Fix the `InvalidCryptoBackendError` exception triggering a `TypeError`. Patch from dz, fixes #13.

4.6 Release 1.4 (7/24/2009)

- Fix bug with `hmac` on Python 2.4. Patch from toshio, closes ticket #2133 from the TurboGears2 Trac.
- Fix bug with occasional `ValueError` from `FileNamespaceManager.do_open`. Fixes #10.
- Fixed bug with session files being saved despite being new and not saved.
- Fixed bug with `CacheMiddleware` overwriting configuration with default arguments despite prior setting.
- Fixed bug with `SyntaxError` not being caught properly in entry point discovery.
- Changed to using `BlobProperty` for Google Datastore.
- Added domain/path properties to the session. This allows one to dynamically set the cookie's domain and/or path on the fly, which will then be set on the cookie for the session.
- Added support for cookie-based sessions in Jython via the JCE (Java Cryptography Extensions). Patch from Alex Grönholm.
- Update Beaker database extensions to work with SQLAlchemy 0.6 PostgreSQL, and Jython.

4.7 Release 1.3.1 (5/5/2009)

- Added a whole bunch of Sphinx documentation for the updated site.
- Added corresponding `remove` as an alias to the `remove_value`.
- Fixed cookie session not having an `invalidate` function.
- Fix bug with `CacheMiddleware` not using proper function to load configuration options, missing the cache regions.

4.8 Release 1.3 (4/6/2009)

- Added last_accessed attribute to session to indicate the previous time the session was last accessed.
- Added setuptools entry points to dynamically discover additional namespace backends.
- Fixed bug with invalidate and locks, fixes #594.
- Added cache.cache decorator for arbitrary caching.
- Added cache.region decorator to the CacheManager object.
- Added cache regions. Can be provided in a configuration INI type, or by adding in a cache_regions arg to the CacheManager.
- Fix bug with timeout not being saved properly.
- Fix bug with cookie-only sessions sending cookies for new sessions even if they weren't supposed to be saved.
- Fix bug that caused a non-auto accessed session to not record the time it was previously accessed resulting in session timeouts.
- Add function to parse configuration dicts as appropriate for use with the CacheManager.
- The "expiretime" is no longer passed to the memcached backend - since if memcached makes the expired item unavailable at the same time the container expires it, then all actors must block until the new value is available (i.e. breaks the anti-dogpile logic).

4.9 Release 1.2.3 (3/2/2009)

- Fix accessed increment to take place *after* the accessed time is checked to see if it has expired. Fixes #580.
- data_dir/lock_dir parameters are optional to most backends; if not present, mutex-based locking will be used for creation functions
- Adjustments to Container to better account for backends which don't provide read/write locks, such as memcached. As a result, the plain "memory" cache no longer requires read/write mutexing.

4.10 Release 1.2.2 (2/14/2009)

- Fix delete bug reported by andres with session not being deleted.

4.11 Release 1.2.1 (2/09/2009)

- Fix memcached behavior as memcached returns None on nonexistent key fetch which broke invalid session checking.

4.12 Release 1.2 (1/22/2009)

- Updated session to only save to the storage *once* no under any/all conditions rather than every time `save()` is called.
- Added `session.revert()` function that reverts the session to the state at the beginning of the request.
- Updated session to store entire session data in a single namespace key, this lets memcached work properly, and makes for more efficient use of the storage system for sessions.

4.13 Release 1.1.3 (12/29/2008)

- Fix the 1.1.2 old cache/session upgrader to handle the `has_current_value` method.
- Make `InvalidCacheBackendError` an `ImportError`.

4.14 Release 1.1.2 (11/24/2008)

- Upgrade Beaker pre-1.1 cache/session values to the new format rather than throwing an exception.

4.15 Release 1.1.1 (11/24/2008)

- Fixed bug in Google extension which passed arguments it should no longer pass to `NamespaceManager`.
- Fixed bug involving lockfiles left open during cache “value creation” step.

4.16 Release 1.1 (11/16/2008)

- file-based cache will not hold onto cached value once read from file; will create new value if the file is deleted as opposed to re-using what was last read. This allows external removal of files to be used as a cache-invalidation mechanism.
- file-based locking will not unlink lockfiles; this can interfere with the `flock()` mechanism in the event that a concurrent process is accessing the files.
- Sending “type” and other namespace config arguments to `cache.get()/cache.put()/cache.remove_value()` is deprecated. The namespace configuration is now preferred at the Cache level, i.e. when you construct a `Cache` or call `cache_manager.get_cache()`. This removes the ambiguity of `Cache`’s dictionary interface and `has_key()` methods, which have no awareness of those arguments.
- the “expiretime” in use is stored in the cache itself, so that it is always available when calling `has_key()` and other methods. Between this change and the deprecation of ‘type’, the `Cache` no longer has any need to store cache configuration in memory per cache key, which in a dynamically-generated key scenario stores an arbitrarily large number of configurations - essentially a memory leak.
- memcache caching has been vastly improved, no longer stores a list of all keys, which along the same theme prevented efficient usage for an arbitrarily large number of keys. The `keys()` method is now unimplemented, and `cache.remove()` clears the entire memcache cache across all namespaces. This is what the memcache API provides so it’s the best we can do.

- memcache caching passes along “expiretime” to the memcached “time” parameter, so that the cache itself can reduce its size for elements which are expired (memcache seems to manage its size in any case, this is just a hint to improve its operation).
- replaced homegrown ThreadLocal implementation with threading.local, falls back to a 2.3 compat one for python<2.4

4.17 Release 1.0.3 (10/14/2008)

- Fixed os.getpid issue on GAE.
- CookieSession will add ‘_expires’ value to data when an expire time is set, and uses it

4.18 Release 1.0.2 (9/22/2008)

- Fixed bug caused when attempting to invalidate a session that hadn’t previously been created.

4.19 Release 1.0.1 (8/19/2008)

- Bug fix for cookie sessions to retain id before clearing values.

4.20 Release 1.0 (8/13/2008)

- Added cookie delete to both cookie only sessions and normal sessions, to help with proxies and such that may determine whether a user is logged in via a cookie. (cookie varies, etc.). Suggested by Felix Schwarz.
- cache.get_value() now uses the given ****kwargs** in all cases in the same manner as cache.set_value(). This way you can send a new createfunc to cache.get_value() each time and it will be used.

4.21 Release 0.9.5 (6/19/2008)

- Fixed bug in memcached to be tolerant of keys disappearing when memcached expires them.
- Fixed the cache functionality to actually work, previously set_value was ignored if there was already a value set.

4.22 Release 0.9.4 (4/13/2008)

- Adding ‘google’ backend datastore, available by specifying ‘google’ as the cache/session type. Note that this takes an optional table_name used to name the model class used.
- SECURITY BUG: Fixed security issue with Beaker not properly removing directory escaping characters from the session ID when un-signed sessions are used. Reported with patch by Felix Schwarz.
- Fixed bug with Beaker not playing well with Registry when its placed above it in the stack. Thanks Wichert Akkerman.

4.23 Release 0.9.3 (2/28/2008)

- Adding 'id' to cookie-based sessions for better compatibility.
- Fixed error with exception still raised for PyCrypto missing.
- WARNING: Session middleware no longer catches Paste HTTP Exceptions, apps are now expected to capture and handle Paste HTTP Exceptions themselves.
- Fixed Python 2.4 compatibility bug in hmac.
- Fixed key lookup bug on cache object to only use the settings for the key lookup. Found by Andrew Stromnov.

4.24 Release 0.9.2 (2/13/2008)

- Added option to make Beaker use a secure cookie.
- Removed CTRCipher as pycryptopp doesn't need it.
- Changed AES to use 256 bit.
- Fixed signing code to use hmac with sha for better signing security.
- Fixed memcached code to use delete_multi on clearing the keys for efficiency and updated key retrieval to properly store and retrieve None values.
- Removing cookie.py and signed cookie middleware, as the environ_key option for session middleware provides a close enough setting.
- Added option to use just cookie-based sessions without requiring encryption.
- Switched encryption requirement from PyCrypto to pycryptopp which uses a proper AES in Counter Mode.

4.25 Release 0.9.1 (2/4/2008)

- Fixed bug in middleware using module that wasn't imported.

4.26 Release 0.9 (12/17/07)

- Fixed bug in memcached replace to actually replace spaces properly.
- Fixed md5 cookie signature to use SHA-1 when available.
- Updated cookie-based session storage to use 256-bit AES-CTR mode with a SHA-1 HMAC signature. Now requires PyCrypto to use for AES scheme.
- WARNING: Moved session and cache middleware to middleware, as per the old deprecation warnings had said was going to happen for 0.8.
- Added cookie-only session storage with RC4 ciphered encryption, requires Python 2.4.
- Add the ability to specify the cookie's domain for sessions.

4.27 Release 0.8.1 (11/15/07)

- Fixed bug in database.py not properly handling missing sqlalchemy library.

4.28 Release 0.8 (10/17/07)

- Fixed bug in prior db update causing session to occasionally not be written back to the db.
- Fixed memcached key error with keys containing spaces. Thanks Jim Musil.
- WARNING: Major change to ext:database to use a single row per namespace. Additionally, there's an accessed and created column present to support easier deletion of old cache/session data. You *will* need to drop any existing tables being used by the ext:database backend.
- Streamline ext:database backend to avoid unnecessary database selects for repeat data.
- Added SQLAlchemy 0.4 support to ext:database backend.

4.29 Release 0.7.5 (08/18/07)

- Fixed data_dir parsing for session string coercions, no longer picks up None as a data_dir.
- Fixed session.get_by_id to lookup recently saved sessions properly, also updates session with creation/access time upon save.
- Add unit tests for get_by_id function. Updated get_by_id to not result in additional session files.
- Added session.get_by_id function to retrieve a session of the given id.

4.30 Release 0.7.4 (07/09/07)

- Fixed issue with Beaker not properly handling arguments as Pylons may pass them in.
- Fixed unit test to catch file removal exception.
- Fixed another bug in synchronization, this one involving reentrant conditions with file synchronization
- If a file open fails due to pickling errors, locks just opened are released unconditionally

4.31 Release 0.7.3 (06/08/07)

- Beaker was not properly parsing input options to session middleware. Thanks to Yannick Gingras and Timothy S for spotting the issue.
- Changed session to only send the cookie header if its a new session and save() was called. Also only creates the session file under these conditions.

4.32 Release 0.7.2 (05/19/07)

- Added deprecation warning for middleware move, relocated middleware to cache and session modules for backwards compatibility.

4.33 Release 0.7.1 05/18/07)

- adjusted synchronization logic to account for Mako/new Cache object's multithreaded usage of Container.

4.34 Release 0.7 (05/18/07)

- WARNING: Cleaned up Cache object based on Mako cache object, this changes the call interface slightly for creating a Cache object directly. The middleware cache object is unaffected from an end-user view. This change also avoids duplicate creations of Cache objects.
- Adding database backend and unit tests.
- Added memcached test, fixed memcached namespace arg passing.
- Fixed session and cache tests, still failing syncdict test. Added doctests for Cache and Session middleware.
- Cleanup of container/cache/container_test
- Namespaces no longer require a context, removed NamespaceContext?
- Logging in container.py uses logging module
- Cleanup of argument passing, use name ****kwargs** instead of ****params** for generic kwargs
- Container classes contain a static create_namespace() method, namespaces are accessed from the ContainerContext? via string name + container class alone
- Implemented (but not yet tested) clear() method on Cache, locates all Namespaces used thus far and clears each one based on its keys() collection
- Fixed Cache.clear() method to actually clear the Cache namespace.
- Updated memcached backend to split servers on ';' for multiple memcached backends.
- Merging MyghtyUtils code into Beaker.

4.35 Release 0.6.3 (03/18/2007)

- Added api with customized Session that doesn't require a Myghty request object, just a dict. Updated session to use the new version.
- Removing unicode keys as some dbm backends can't handle unicode keys.
- Adding core files that should've been here.
- More stringent checking for existence of a session.
- Avoid recreating the session object when it's empty.

INDICES AND TABLES

- *Index*
- *Module Index*
- *Search Page*
- *Glossary*

5.1 Module Listing

5.1.1 `beaker.cache` – Cache module

Cache object

The Cache object is used to manage a set of cache files and their associated backend. The backends can be rotated on the fly by specifying an alternate type when used.

Advanced users can add new backends in `beaker.backends`

Module Contents

`cache_region` (*region*, **deco_args*)

Decorate a function to cache itself using a cache region

The region decorator requires arguments if there are more than 2 of the same named function, in the same module. This is because the namespace used for the functions cache is based on the functions name and the module.

Example:

```
# Add cache region settings to beaker:
beaker.cache.cache_regions.update(dict_of_config_region_options))

@cache_region('short_term', 'some_data')
def populate_things(search_term, limit, offset):
    return load_the_data(search_term, limit, offset)

return load('rabbits', 20, 0)
```

Note: The function being decorated must only be called with positional arguments.

region_invalidate (*namespace, region, *args*)

Invalidate a cache region namespace or decorated function

This function only invalidates cache spaces created with the `cache_region` decorator.

Parameters

- *namespace* – Either the namespace of the result to invalidate, or the cached function reference
- *region* – The region the function was cached to. If the function was cached to a single region then this argument can be `None`
- *args* – Arguments that were used to differentiate the cached function as well as the arguments passed to the decorated function

Example:

```
# Add cache region settings to beaker:
beaker.cache.cache_regions.update(dict_of_config_region_options))

def populate_things(invalidate=False):

    @cache_region('short_term', 'some_data')
    def load(search_term, limit, offset):
        return load_the_data(search_term, limit, offset)

    # If the results should be invalidated first
    if invalidate:
        region_invalidate(load, None, 'some_data',
                           'rabbits', 20, 0)

    return load('rabbits', 20, 0)
```

class Cache (*namespace, type='memory', expiretime=None, starttime=None, expire=None, **nsargs*)

Front-end to the containment API implementing a data cache.

Parameters

- *namespace* – the namespace of this Cache
- *type* – type of cache to use
- *expire* – seconds to keep cached data
- *expiretime* – seconds to keep cached data (legacy support)
- *starttime* – time when cache was cache was

get (*key, **kw*)

Retrieve a cached value from the container

clear ()

Clear all the values from the namespace

class CacheManager (***kwargs*)

Initialize a CacheManager object with a set of options

Options should be parsed with the `parse_cache_config_options()` function to ensure only valid options are used.

region (*region, *args*)

Decorate a function to cache itself using a cache region

The region decorator requires arguments if there are more than 2 of the same named function, in the same module. This is because the namespace used for the functions cache is based on the functions name and the module.

Example:

```
# Assuming a cache object is available like:
cache = CacheManager(dict_of_config_options)

def populate_things():

    @cache.region(' short_term ', ' some_data ')
    def load(search_term, limit, offset):
        return load_the_data(search_term, limit, offset)

    return load(' rabbits ', 20, 0)
```

Note: The function being decorated must only be called with positional arguments.

region_invalidate (*namespace, region, *args*)

Invalidate a cache region namespace or decorated function

This function only invalidates cache spaces created with the cache_region decorator.

Parameters

- *namespace* – Either the namespace of the result to invalidate, or the name of the cached function
- *region* – The region the function was cached to. If the function was cached to a single region then this argument can be None
- *args* – Arguments that were used to differentiate the cached function as well as the arguments passed to the decorated function

Example:

```
# Assuming a cache object is available like:
cache = CacheManager(dict_of_config_options)

def populate_things(invalidate=False):

    @cache.region(' short_term ', ' some_data ')
    def load(search_term, limit, offset):
        return load_the_data(search_term, limit, offset)

    # If the results should be invalidated first
    if invalidate:
        cache.region_invalidate(load, None, ' some_data ',
                                ' rabbits ', 20, 0)

    return load(' rabbits ', 20, 0)
```

cache (**args, **kwargs*)

Decorate a function to cache itself with supplied parameters

Parameters

- *args* – Used to make the key unique for this function, as in region() above.
- *kwargs* – Parameters to be passed to get_cache(), will override defaults

Example:

```
# Assuming a cache object is available like:
cache = CacheManager(dict_of_config_options)

def populate_things():

    @cache.cache(' mycache ', expire=15)
    def load(search_term, limit, offset):
        return load_the_data(search_term, limit, offset)

    return load(' rabbits ', 20, 0)
```

Note: The function being decorated must only be called with positional arguments.

invalidate (*func*, **args*, ***kwargs*)

Invalidate a cache decorated function

This function only invalidates cache spaces created with the cache decorator.

Parameters

- *func* – Decorated function to invalidate
- *args* – Used to make the key unique for this function, as in region() above.
- *kwargs* – Parameters that were passed for use by get_cache(), note that this is only required if a type was specified for the function

Example:

```
# Assuming a cache object is available like:
cache = CacheManager(dict_of_config_options)

def populate_things(invalidate=False):

    @cache.cache(' mycache ', type=" file ", expire=15)
    def load(search_term, limit, offset):
        return load_the_data(search_term, limit, offset)

    # If the results should be invalidated first
    if invalidate:
        cache.invalidate(load, ' mycache ', ' rabbits ', 20, 0, type=" file ")
    return load(' rabbits ', 20, 0)
```

5.1.2 beaker.container – Container and Namespace classes

Container and Namespace classes

Module Contents

class Container()

class ContainerMeta (classname, bases, dict_)

class DBMContainer()

```
class DBMNamespaceManager (namespace, dbmmodule=None, data_dir=None, dbm_dir=None, lock_dir=None, digest_filenames=True, **kwargs)
```

```
class FileContainer ()
```

```
class FileNamespaceManager (namespace, data_dir=None, file_dir=None, lock_dir=None, digest_filenames=True, **kwargs)
```

```
class MemoryContainer ()
```

```
class MemoryNamespaceManager (namespace, **kwargs)
```

```
class NamespaceManager (namespace)
```

Handles dictionary operations and locking for a namespace of values.

The implementation for setting and retrieving the namespace data is handled by subclasses.

NamespaceManager may be used alone, or may be privately accessed by one or more Container objects. Container objects provide per-key services like expiration times and automatic recreation of values.

Multiple NamespaceManagers created with a particular name will all share access to the same underlying datasource and will attempt to synchronize against a common mutex object. The scope of this sharing may be within a single process or across multiple processes, depending on the type of NamespaceManager used.

The NamespaceManager itself is generally threadsafe, except in the case of the DBMNamespaceManager in conjunction with the gdbm dbm implementation.

```
set_value (key, value, expiretime=None)
```

Optional set_value() method called by Value.

Allows an expiretime to be passed, for namespace implementations which can prune their collections using expiretime.

```
class OpenResourceNamespaceManager (namespace)
```

A NamespaceManager where read/write operations require opening/ closing of a resource which is possibly mutexed.

```
class Value (key, namespace, createfunc=None, expiretime=None, starttime=None)
```

```
has_value ()
```

return true if the container has a value stored.

This is regardless of it being expired or not.

5.1.3 beaker.middleware – Middleware classes

Module Contents

```
class CacheMiddleware (app, config=None, environ_key='beaker.cache', **kwargs)
```

Initialize the Cache Middleware

The Cache middleware will make a Cache instance available every request under the `environ['beaker.cache']` key by default. The location in environ can be changed by setting `environ_key`.

config dict All settings should be prefixed by 'cache.'. This method of passing variables is intended for Paste and other setups that accumulate multiple component settings in a single dictionary. If config contains *no cache. prefixed args*, then *all* of the config options will be used to initialize the Cache objects.

environ_key Location where the Cache instance will keyed in the WSGI environ

****kwargs** All keyword arguments are assumed to be cache settings and will override any settings found in `config`

class SessionMiddleware (*wrap_app, config=None, environ_key='beaker.session', **kwargs*)
Initialize the Session Middleware

The Session middleware will make a lazy session instance available every request under the `environ['beaker.session']` key by default. The location in environ can be changed by setting `environ_key`.

config dict All settings should be prefixed by 'session.'. This method of passing variables is intended for Paste and other setups that accumulate multiple component settings in a single dictionary. If `config` contains *no cache. prefixed args*, then *all* of the `config` options will be used to initialize the Cache objects.

environ_key Location where the Session instance will keyed in the WSGI environ

****kwargs** All keyword arguments are assumed to be session settings and will override any settings found in `config`

5.1.4 beaker.session – Session classes

Module Contents

class CookieSession (*request, key='beaker.session.id', timeout=None, cookie_expires=True, cookie_domain=None, encrypt_key=None, validate_key=None, secure=False, **kwargs*)
Pure cookie-based session

Options recognized when using cookie-based sessions are slightly more restricted than general sessions.

key The name the cookie should be set to.

timeout How long session data is considered valid. This is used regardless of the cookie being present or not to determine whether session data is still valid.

encrypt_key The key to use for the session encryption, if not provided the session will not be encrypted.

validate_key The key used to sign the encrypted session

cookie_domain Domain to use for the cookie.

secure Whether or not the cookie should only be sent over SSL.

save (*accessed_only=False*)
Saves the data for this session to persistent storage

expire ()
Delete the 'expires' attribute on this Session, if any.

delete ()
Delete the cookie, and clear the session

invalidate ()
Clear the contents and start a new session

```
class Session (request, id=None, invalidate_corrupt=False, use_cookies=True, type=None, data_dir=None,  

               key='beaker.session.id', timeout=None, cookie_expires=True, cookie_domain=None, secret=None,  

               secure=False, namespace_class=None, **namespace_args)
```

Session object that uses container package for storage.

key The name the cookie should be set to.

timeout How long session data is considered valid. This is used regardless of the cookie being present or not to determine whether session data is still valid.

cookie_domain Domain to use for the cookie.

secure Whether or not the cookie should only be sent over SSL.

save (*accessed_only=False*)

Saves the data for this session to persistent storage

If *accessed_only* is True, then only the original data loaded at the beginning of the request will be saved, with the updated last accessed time.

revert ()

Revert the session to its original state from its first access in the request

lock ()

Locks this session against other processes/threads. This is automatic when load/save is called.

use with caution and always with a corresponding 'unlock' inside a "finally:" block, as a stray lock typically cannot be unlocked without shutting down the whole application.

unlock ()

Unlocks this session against other processes/threads. This is automatic when load/save is called.

use with caution and always within a "finally:" block, as a stray lock typically cannot be unlocked without shutting down the whole application.

delete ()

Deletes the session from the persistent storage, and sends an expired cookie out

invalidate ()

Invalidates this session, creates a new session id, returns to the *is_new* state

```
class SessionObject (environ, **params)
```

Session proxy/lazy creator

This object proxies access to the actual session object, so that in the case that the session hasn't been used before, it will be setup. This avoid creating and loading the session from persistent storage unless its actually used during the request.

persist ()

Persist the session to the storage

If its set to autosave, then the entire session will be saved regardless of if *save()* has been called. Otherwise, just the accessed time will be updated if *save()* was not called, or the session will be saved if *save()* was called.

get_by_id (*id*)

Loads a session given a session ID

accessed ()

Returns whether or not the session has been accessed

```
class SignedCookie (secret, input=None)
```

Extends python cookie to give digital signature support

5.1.5 beaker.synchronization – Synchronization classes

Synchronization functions.

File- and mutex-based mutual exclusion synchronizers are provided, as well as a name-based mutex which locks within an application based on a string name.

Module Contents

class ConditionSynchronizer (*identifier*)
a synchronizer using a Condition.

class FileSynchronizer (*identifier, lock_dir*)
a synchronizer which locks using flock().

Adapted for Python/multithreads from Apache::Session::Lock::File,
<http://search.cpan.org/src/CWEST/Apache-Session-1.81/Session/Lock/File.pm>

This module does not unlink temporary files, because it interferes with proper locking. This can cause problems on certain systems (Linux) whose file systems (ext2) do not perform well with lots of files in one directory. To prevent this you should use a script to clean out old files from your lock directory.

class NameLock (*identifier=None, reentrant=False*)
a proxy for an RLock object that is stored in a name based registry.

Multiple threads can get a reference to the same RLock based on the name alone, and synchronize operations related to that name.

class SynchronizerImpl ()

5.1.6 beaker.util – Beaker Utilities

Beaker utilities

Module Contents

encoded_path (*root, identifiers, extension='enc', depth=3, digest_filenames=True*)
Generate a unique file-accessible path from the given list of identifiers starting at the given root directory.

func_namespace (*func*)
Generates a unique namespace for a function

class SyncDict ()
An efficient/threadsafe singleton map algorithm, a.k.a. “get a value based on this key, and create if not found or not valid” paradigm:

exists && isvalid ? get : create

Designed to work with weakref dictionaries to expect items to asynchronously disappear from the dictionary.

Use python 2.3.3 or greater ! a major bug was just fixed in Nov. 2003 that was driving me nuts with garbage collection/weakrefs in this section.

class ThreadLocal ()
stores a value on a per-thread basis

verify_directory (*dir*)

verifies and creates a directory. tries to ignore collisions with other threads and processes.

parse_cache_config_options (*config*, *include_defaults=True*)

Parse configuration options and validate for use with the CacheManager

5.1.7 beaker.ext.database – Database Container and NameSpace Manager classes

Module Contents

class DatabaseContainer ()

class DatabaseNamespaceManager (*namespace*, *url=None*, *sa_opts=None*, *optimistic=False*, *table_name='beaker_cache'*, *data_dir=None*, *lock_dir=None*, ***params*)

Creates a database namespace manager

url SQLAlchemy compliant db url

sa_opts A dictionary of SQLAlchemy keyword options to initialize the engine with.

optimistic Use optimistic session locking, note that this will result in an additional select when updating a cache value to compare version numbers.

table_name The table name to use in the database for the cache.

5.1.8 beaker.ext.google – Google Container and NameSpace Manager classes

Module Contents

class GoogleContainer ()

class GoogleNamespaceManager (*namespace*, *table_name='beaker_cache'*, ***params*)

Creates a datastore namespace manager

5.1.9 beaker.ext.memcached – Memcached Container and NameSpace Manager classes

Module Contents

class MemcachedContainer ()

class MemcachedNamespaceManager (*namespace*, *url=None*, *data_dir=None*, *lock_dir=None*, ***params*)

5.1.10 beaker.ext.sqla – SQLAlchemy Container and NameSpace Manager classes

Module Contents

make_cache_table (*metadata*, *table_name='beaker_cache'*)

Return a Table object suitable for storing cached values for the namespace manager. Do not create the table.

class SqlaContainer ()

class SqlaNamespaceManager (*namespace, bind, table, data_dir=None, lock_dir=None, **kwargs*)

Create a namespace manager for use with a database table via SQLAlchemy.

bind SQLAlchemy Engine or Connection object

table SQLAlchemy Table object in which to store namespace data. This should usually be something created by `make_cache_table`.

5.1.11 **beaker.crypto.pbkdf2** – PKCS#5 v2.0 Password-Based Key Derivation classes

Module Contents

crypt (*word, salt=None, iterations=None*)

PBKDF2-based unix crypt(3) replacement.

The number of iterations specified in the salt overrides the 'iterations' parameter.

The effective hash length is 192 bits.

```
class PBKDF2 (passphrase, salt, iterations=1000, digestmodule=<module 'Crypto.Hash.SHA'
from
'/Users/ben/Programming/Python/py2.6devenv/lib/python2.6/site-
packages/Crypto/Hash/SHA.pyc'>, macmodule=<module 'Crypto.Hash.HMAC'
from
'/Users/ben/Programming/Python/py2.6devenv/lib/python2.6/site-
packages/Crypto/Hash/HMAC.pyc'>)
```

PBKDF2.py : PKCS#5 v2.0 Password-Based Key Derivation

This implementation takes a passphrase and a salt (and optionally an iteration count, a digest module, and a MAC module) and provides a file-like object from which an arbitrarily-sized key can be read.

If the passphrase and/or salt are unicode objects, they are encoded as UTF-8 before they are processed.

The idea behind PBKDF2 is to derive a cryptographic key from a passphrase and a salt.

PBKDF2 may also be used as a strong salted password hash. The 'crypt' function is provided for that purpose.

Remember: Keys generated using PBKDF2 are only as strong as the passphrases they are derived from.

close()

Close the stream.

hexread (*octets*)

Read the specified number of octets. Return them as hexadecimal.

Note that `len(obj.hexread(n)) == 2*n`.

read (*bytes*)

Read the specified number of key bytes.